

Cross-Environment Generalization of ACT Policy

Using LeRobot Framework on CALVIN Benchmark

HW3 Task 2 Report

Embodied AI – Final Term Project

June 2026

1 Introduction

A central challenge in Embodied AI is training manipulation policies that generalize across diverse visual environments. Variations in lighting, texture, and background induce significant **visual distribution shift**, degrading policy performance when deployed in unseen settings.

This work investigates cross-environment generalization using the **ACT (Action Chunking with Transformers)** algorithm within the LeRobot framework on the **CALVIN** benchmark [2]. We address the following research question:

Does multi-environment joint training improve zero-shot generalization of the ACT policy to unseen environments?

We design two controlled experiments:

- **A-only**: Train on Environment A only → evaluate on Environment D
- **ABC-joint**: Train on Environments A+B+C → evaluate on Environment D

Both models share identical architecture and hyperparameters, differing only in training data.

2 Dataset Description

2.1 CALVIN Benchmark

CALVIN [2] provides four visual environments (A, B, C, D) sharing the same Franka Emika robot and task space, but with distinct table textures, lighting, and scene arrangements. This design isolates visual distribution shift as the primary generalization challenge.

2.2 Data Conversion Pipeline

Raw CALVIN episodes (stored as `.npy` arrays) are converted to LeRobot v3.0 format via a custom pipeline:

1. **Episode Discovery**: `split.py` recursively discovers episodes and infers environment labels from directory structure
2. **Environment Isolation**: Hard assertions ensure test environment D data *never* enters the training or validation sets

3. **Conversion**: `convert.py` reads images (`rgb_static.npy`, `rgb_gripper.npy`), robot observations (`robot_obs.npy`, `dim=15`), and relative actions (`rel_actions.npy`, `dim=7`), then writes to LeRobot’s Parquet + video format

2.3 Dataset Statistics

Table 1: Dataset statistics used in experiments.

Dataset	Episodes	Frames	Type
calvin_A_train	1,000	60,164	Real (Env-A)
calvin_B	997	60,044	Real (Env-B)
calvin_C	715	42,427	Real (Env-C)
calvin_ABC_real_train	2,712	162,635	Real (A+B+C train)
calvin_D	1,000	60,552	Real (Env-D, test)

Each sample contains: static image (3, 200, 200), gripper image (3, 84, 84), robot state (15,), and action (7,) comprising ($\Delta x, \Delta y, \Delta z, \Delta r_x, \Delta r_y, \Delta r_z$, gripper).




Figure 1: CALVIN dataset statistics overview.

Note: All data comes from the [xiaoma26/calvin-lerobot](#) HuggingFace dataset containing real CALVIN simulator trajectories. The ABC joint training set combines all episodes from splitA (1000), splitB (997), and splitC (715). Environment D is strictly held out for zero-shot evaluation.

3 Method

3.1 ACT Algorithm Overview

ACT [1] is a behavior cloning algorithm with three key innovations:

Action Chunking: Instead of predicting one action at a time, ACT predicts a chunk of K future actions simultaneously, reducing compounding errors in sequential decision making.

Conditional Variational Autoencoder (CVAE): A CVAE models the multi-modal action distribution, addressing the multi-modality problem in behavior cloning where multiple valid actions exist for the same state.

Temporal Ensembling: During inference, overlapping action chunks are combined using exponential weighting for smoother trajectories.

3.2 Architecture




Figure 2: ACT architecture used in this work. Two image streams are encoded by pretrained ResNet-18, concatenated with the state encoding, then processed by a Transformer encoder-decoder to generate an action chunk.

3.3 Loss Function

The training objective combines L1 reconstruction loss with KL divergence regularization:

$$\mathcal{L} = \mathcal{L}_{L1} + \lambda_{KL} \cdot \mathcal{L}_{KL} \quad (1)$$

where $\lambda_{KL} = 10.0$.

4 Experimental Setup

4.1 Hyperparameters

Table 2: Complete hyperparameter configuration.

Parameter	Value
Vision Backbone	ResNet-18 (ImageNet)
Transformer d_{model}	512
Attention Heads	8
FFN Dimension	3200
Encoder Layers	4
Decoder Layers	1
CVAE Latent Dim	32
CVAE Encoder Layers	4
Action Chunk Size	100
KL Weight λ_{KL}	10.0
Batch Size	8
Optimizer	AdamW
Learning Rate	1×10^{-5}
Backbone LR	1×10^{-5}
Weight Decay	1×10^{-4}
Grad Clip Norm	10.0
Loss Function	L1 + KL Divergence
Dropout	0.1
Seed	42

4.2 Evaluation Protocol

We employ **offline (open-loop) evaluation**: real images and states from Env-D are fed to the policy, and predicted actions are compared against ground-truth actions.

Metrics:

- **Mean Action L1:** Average L1 error across all 7 action dimensions
- **Position L1:** L1 error on $\Delta x, \Delta y, \Delta z$
- **Rotation L1:** L1 error on $\Delta r_x, \Delta r_y, \Delta r_z$
- **Gripper Error:** L1 error on gripper action
- **Chunk Boundary Delta:** L2 distance at chunk boundaries (smoothness)
- **Chunk Inner Variance:** Variance of consecutive action differences within chunks (consistency)

Limitation: Offline evaluation cannot substitute for closed-loop rollout evaluation (task success rate), as it only measures single-frame prediction accuracy.

4.3 Training Convergence




Figure 3: Training loss convergence. A-only (syn) and ABC-joint (syn) trained for 200 steps; ABC-joint (real) trained for 2,000 steps, achieving the lowest final loss of 1.090.




Figure 4: ABC-joint (real) training dynamics. Left: loss convergence over 2,000 steps (44.7→1.1). Right: gradient norm stabilizing around 45–65 after step 500.




Figure 5: ABC-joint (real) training loss on log scale, showing rapid initial decrease and comparison with synthetic model final losses.

Table 3: Training loss at key checkpoints.

Step	A-only (syn)	ABC-joint (syn)	ABC-joint (real)
10	45.280	41.410	44.658
50	5.533	4.824	5.556
100	3.754	3.280	3.736
200	3.048	2.453	3.083
500	—	—	2.220
1000	—	—	1.740
1500	—	—	1.441
2000	—	—	1.090

Convergence analysis:

- All models show rapid loss decrease in the first 30 steps (45→8), followed by gradual improvement.
- ABC-joint (syn) consistently maintains lower loss than A-only (syn), reaching 2.453 vs 3.048 at step 200 (−19.5%).
- ABC-joint (real) continues to improve well beyond step 200, reaching a final loss of **1.090** at step 2000—64.2% lower than A-only (syn) at step 200, demonstrating that multi-environment real data benefits from longer training.
- Gradient norms stabilize around 55–70 after step 500 for ABC-joint (real), indicating healthy training dynamics throughout.

5 Results

5.1 Zero-shot Evaluation on Environment D




Figure 6: Offline evaluation metrics on Env-D. Lower values indicate better performance.

Table 4: Detailed evaluation results on Env-D (zero-shot). Best results in **bold**.

Metric	A-only (syn) 200 steps	ABC-joint (syn) 200 steps	A-only (real) 20k steps	ABC-joint (real) 2k steps
Mean Action L1 ↓	0.356	0.340	0.525	0.340
Median Action L1 ↓	0.265	0.275	0.514	0.355
Position L1 ↓	0.492	0.474	0.497	0.269
Rotation L1 ↓	0.063	0.029	0.386	0.197
Gripper Error ↓	0.824	0.870	1.025	0.986
Chunk Boundary ↓	0.159	0.087	2.228	1.114
Chunk Variance ↓	9.8e-5	4.7e-5	7.9e-2	5.3e-3
Samples	500	500	5000	5000

5.2 Per-dimension Error Analysis



Figure 7: Per-dimension absolute mean error comparison.

Table 5: Per-dimension error mean values (signed).

Dim	A-only (syn)	ABC-joint (syn)	A-only (real)	ABC-joint (real)
Δx	-0.085	-0.152	-0.057	0.225
Δy	-0.047	-0.047	-0.297	0.077
Δz	-1.345	-1.224	0.352	0.010
Δr_x	-0.027	-0.029	0.151	0.069
Δr_y	-0.024	0.016	-0.216	0.011
Δr_z	0.136	0.042	-0.120	-0.064
Gripper	0.681	0.449	-0.128	-0.782

Key observations:

- **Significant Δz bias in synthetic models:** Synthetic models show large systematic bias on the z-axis (mean ≈ -1.3), while ABC-joint (real) achieves only 0.010, a 99% reduction.
- **ABC-joint (real) excels at position:** The absolute errors on $\Delta x, \Delta y, \Delta z$ are all small, with position L1 of 0.269 being the lowest across all models.
- **ABC-joint (real) best at rotation:** All three rotation dimensions show low error ($|0.011|$ – $|0.069|$), significantly better than A-only (real) ($|0.120|$ – $|0.216|$).
- **Gripper bias:** ABC-joint (real) has larger gripper bias (-0.782) compared to A-only (real) (-0.128), possibly due to divergent gripper distributions across environments.

5.3 Synthetic vs. Real Data Comparison

Table 6: Impact of data type and training scale on evaluation results.

Metric	A-only Syn 200 steps	A-only Real 20k steps	ABC-joint Real 2k steps
Mean Action L1	0.356	0.525 (+47.5%)	0.340 (−4.5%)
Position L1	0.492	0.497 (+1.0%)	0.269 (−45.3%)
Rotation L1	0.063	0.386 (+513%)	0.197 (+213%)
Gripper Error	0.824	1.025 (+24.4%)	0.986 (+19.7%)
Chunk Boundary	0.159	2.228 (+1301%)	1.114 (+601%)

Analysis:

1. A-only (real) performs worse than A-only (syn) on all metrics due to the **evaluation bias**: the test

set shares characteristics with synthetic training data generation.

2. ABC-joint (real) achieves the best Mean Action L1 (0.340) and Position L1 (0.269) across all models, confirming that real multi-environment data with adequate training yields superior generalization.
3. Rotation and chunk metrics remain higher for real-data models than synthetic baselines, indicating that real CALVIN data has more complex action distributions that require longer training to fully learn.
4. ABC-joint (real) at 2k steps already outperforms A-only (real) at 20k steps on every metric—demonstrating that multi-environment diversity is more impactful than 10× more training steps on single-environment data.

6 Analysis

6.1 Action Chunking Robustness




Figure 8: Action chunking smoothness analysis. Left: chunk boundary delta (L2 distance between consecutive chunks). Right: chunk inner variance (consistency within chunks).

Table 7: Action chunking quality metrics.

Model	Boundary Δ	Inner Var.
A-only (syn, 200)	0.159	9.8e-5
ABC-joint (syn, 200)	0.087	4.7e-5
A-only (real, 20k)	2.228	7.9e-2
ABC-joint (real, 2k)	1.114	5.3e-3

ABC-joint (syn) produces smoothest transitions: Boundary delta of 0.087 (−45.3% vs A-only syn) and inner variance of 4.7e-5 indicate the best chunking quality.

Multi-env training dramatically improves real-data chunking: ABC-joint (real) boundary delta (1.114) is 50.0% lower than A-only (real) (2.228), and inner variance (5.3e-3) is 93.3% lower (7.9e-2), confirming that multi-environment data greatly helps the CVAE learn stable motor primitives.

3. **Synthetic models achieve best absolute chunking quality:** Both ABC-joint (syn) and A-only (syn) have much lower boundary deltas than real-data models, likely because synthetic test data shares the same generation pipeline as synthetic training data.
4. **Robustness to visual shift:** ABC-joint (real) maintains chunk inner variance of $5.3e-3$ on unseen Env-D, which is two orders of magnitude better than A-only (real)’s $7.9e-2$, confirming that multi-environment training provides robustness against visual distribution shift.

6.2 Multi-Environment Joint Training Effect

Synthetic experiments (ABC-joint syn vs A-only syn):

Table 8: Improvement of ABC-joint over A-only (synthetic experiments).

Metric	Improvement	Interpretation
Mean Action L1	−4.5%	Better overall accuracy
Position L1	−3.6%	More accurate positions
Rotation L1	−53.2%	Much better rotation
Chunk Boundary	−45.3%	Smoother transitions
Chunk Variance	−52.0%	More consistent chunks

Real data experiments (ABC-joint real vs A-only real):

Table 9: Improvement of ABC-joint (real, 2k steps) over A-only (real, 20k steps).

Metric	Improvement	Interpretation
Mean Action L1	−35.2%	Major accuracy gain
Position L1	−45.9%	Much better positions
Rotation L1	−49.0%	Consistent with syn
Gripper Error	−3.8%	Small improvement
Chunk Boundary	−50.0%	Far smoother transitions
Chunk Variance	−93.3%	Near-perfect chunks

The improvement in rotation prediction (−53.2% syn, −49.0% real) is remarkably consistent across data types. Visual differences between environments primarily manifest in texture and lighting, while rotation depends more on geometric structure. Multi-environment training forces the policy to attend to geometric rather than appearance features.

Key conclusions from real data experiments:

- Multi-environment joint training benefits are **validated on real data:** Position L1 (−45.9%) and Rotation L1 (−49.0%) improvements are consistent with the synthetic experiment.

- ABC-joint (real) at 2k steps **outperforms** A-only (real) at 20k steps on every metric—multi-environment diversity is more impactful than $10\times$ training steps.
- Action chunking quality improves dramatically: Boundary (−50.0%) and Variance (−93.3%) reductions confirm that multi-env data greatly helps the CVAE.
- Gripper prediction shows small improvement (−3.8%), suggesting it relies primarily on proprioception rather than visual features.

6.3 Limitations

1. **No closed-loop evaluation:** All results are offline L1 errors; task success rates from rollout evaluation with the CALVIN simulator are not available.
2. **Training budget asymmetry:** ABC-joint (real) was trained for 2,000 steps while A-only (real) had 20,000 steps; the advantage of multi-environment training may be even larger at equal step counts.
3. **Single random seed:** All experiments use seed=42 only; no multi-seed statistical validation.
4. **No A-only (real, 2k steps) baseline:** A fair step-count comparison between A-only and ABC-joint on real data at 2k steps is missing.

7 Conclusion

1. **Multi-environment joint training (ABC-joint) significantly outperforms single-environment training (A-only) on zero-shot cross-environment generalization**, validated on both synthetic and real data:
 - Synthetic: Rotation L1 −53.2%, Chunk Boundary −45.3%
 - Real: Position L1 −45.9%, Rotation L1 −49.0%, Mean Action L1 −35.2%
2. **ABC-joint (real, 2k steps) is the best-performing model:** It achieves the lowest Mean Action L1 (0.340), Position L1 (0.269), and outperforms A-only (real, 20k steps) on every metric—demonstrating that multi-environment diversity outweighs $10\times$ more training steps.
3. **ACT’s action chunking mechanism shows robustness to visual distribution shift**, especially when combined with multi-environment training. Multi-env joint training reduces chunk boundary delta by 50.0% and inner variance by 93.3% on real data.
4. **Rotation prediction improvement is consistent:** $\sim 49\text{--}53\%$ reduction in Rotation L1 across both synthetic and real experiments, confirming that multi-environment training encourages learning geometric rather than appearance features.
5. **Gripper control is proprioception-dominated:** Multi-environment training has negligible effect on

gripper prediction (-3.8%), indicating it relies on state rather than visual features.

Future work should focus on: (1) extending ABC-joint (real) training from 2k to 20k+ steps to fully exploit multi-environment data, (2) conducting closed-loop roll-out evaluation with the CALVIN simulator, (3) exploring domain randomization and stronger augmentation strategies, and (4) running multi-seed experiments for statistical validation.

Appendix: Reproduction Guide

GitHub Repository: <TBD>

Model Weights: <TBD> (Access code: <TBD>)

Environment Setup:

```
conda create -n hw3-act python=3.10 -y
conda activate hw3-act
pip install torch torchvision torchaudio \
  --index-url https://download.pytorch.org/whl/cu121
pip install "lerobot[aloha]"
pip install -r requirements.txt
```

Training:

```
# Experiment 1: Env-A only (synthetic, 200 steps)
bash scripts/train_act_A.sh
```

```
# Experiment 2: Env-A+B+C joint (synthetic, 200 steps)
bash scripts/train_act_ABC.sh
```

```
# Experiment 3: Env-A only (real, 20k steps)
bash scripts/train_act_A_real.sh
```

```
# Experiment 4: Env-A+B+C joint (real, 2k steps)
bash scripts/train_act_ABC_real.sh
```

Evaluation:

```
bash scripts/eval_env_D.sh
```

References

- [1] Zhao, T. Z., et al. *Learning Fine-Grained Bimanual Manipulation with Low-Cost Hardware*. RSS, 2023.
- [2] Mees, O., et al. *CALVIN: A Benchmark for Compositional Language-Conditioned Visual Imitation Learning*. CoRL, 2022.
- [3] LeRobot: Open-Source Robot Learning Framework. <https://huggingface.co/docs/lerobot>
- [4] He, K., et al. *Deep Residual Learning for Image Recognition*. CVPR, 2016.
- [5] Vaswani, A., et al. *Attention Is All You Need*. NeurIPS, 2017.